

Efficient AI Use for Product Work

Use the lightest tool that can do the job well, then add power only when the work earns it.

1. Start with the job

Ask what actually requires judgment. Use rules, queries, schemas, or calculations for deterministic work.

3. Match tool to task

Fast model for routine work. Reasoning model for ambiguity. Agent only for multi-step tool use.

5. Batch similar work

Cluster feedback, tickets, research, or eval cases together when comparison is useful.

7. Bound agents

Give goal, permitted tools, escalation path, success condition, and a stopping rule.

9. Verify with software

AI interprets. Software checks. Humans decide when consequences or ambiguity justify it.

2. Shrink context first

Give the smallest useful context. Remove stale material before reaching for a stronger model.

4. Reuse, don't rediscover

Keep project briefs, decisions, eval cases, and good workflows durable so AI does not relearn them each time.

6. Structure the output

Prefer Decision / Evidence / Unknowns / Risks / Next step over free-form conversation.

8. Escalate gradually

Known signal → likely source → adjacent evidence → broader investigation. Do not start with every tool.

10. Measure real cost

Count latency, retries, review time, cleanup, tool calls, and bad decisions — not just tokens.

BEFORE I USE AI

- ☐ Does this require judgment?
- ☐ What is the cheapest capable model?
- ☐ What makes the AI stop?
- ☐ What is the smallest useful context?
- ☐ Can software verify any part?
- ☐ How will I know the result is good?

Optimize for outcome per dollar, not tokens per dollar.