

Meridian Discovery Questions

— Why Each One Matters

The 5 discovery questions Paige chose for the Meridian case study, with the reasoning that changed product direction.

Context

Meridian is an AI customer success assistant for a support team. The case study documents how discovery uncovered hidden assumptions about what reps actually do vs. what process documentation claims they do.

The key insight: **discovery for AI-assisted work has to be more precise than for typical software projects, because you're often trying to specify human judgment, not a process you can already fully define.**

Question 1: System Data Capability

Question: What systems have the information the AI needs to understand customer context?

Why This Matters: - The initial assumption: "We have Zendesk for tickets, we have customer data somewhere" - The reality: Customer data was fragmented across 4 systems with incomplete APIs - The specificity matters: AI can't reason about context it can't access

Changed thinking from: - "AI should read the customer's account details" (too vague)

To: - "AI can only access Zendesk tickets. It cannot reach the account management system, CRM, or billing history without manual API setup that hasn't been prioritized. This is the constraint. Design around it."

Why it shows judgment: - Not accepting the "we have the data" assumption without verification - Recognizing that "incomplete access" means the AI can't reason about certain customer needs - Decision: scope the AI to ticket-only context, acknowledge that will limit its usefulness for customers asking about billing

Result: Avoided building an assistant that would confidently answer questions about billing/account status based on incomplete data (confidence trap).

Question 2: Did Discovery Cover Rep-Level Habits?

Question: When reps actually troubleshoot, what do they do that's not in the documented process?

Why This Matters: - Documentation says: "Follow these 6 steps" - Actual reps: Follow 3 steps, then skip to step 6 because step 4 is "usually not needed", then call a specialist if that doesn't work - AI training on docs → AI gives wrong advice about troubleshooting sequence

Changed thinking from: - "Use the documented troubleshooting steps" (naive)

To: - "Most reps skip step 4. Some reps have workarounds for common error codes that aren't documented. Interview 3-4 reps before building anything, don't just read the wiki."

Why it shows judgment: - Recognizing that process documentation ≠ actual process - Knowing that AI trained on docs would diverge from what experienced reps actually do - Decision: Require rep feedback loops, not just documentation

Result: Avoided training an assistant on outdated procedures that would confuse reps or give bad suggestions.

Question 3: Billing Access Approval Scope

Question: Who actually has permission to look at customer billing history, and what's the review process?

Why This Matters: - Seemed like: "Support reps can see billing" - Actually: Tier 1 reps can see it for some customers only. Tier 2 reps have full access. There's a compliance review for anything involving financial data. PCI and SOX apply. - If the AI generates responses about billing: those responses need human review. That changes the approval workflow.

Changed thinking from: - "AI can answer billing questions if it reads the system" (oversimplified)

To: - "AI can see billing data, but generated responses about it need human approval. This is a Tier 2 review gate. If you want Tier 1 reps to use the AI for billing questions, you need a human-in-the-loop workflow where the AI drafts and a Tier 2 rep approves."

Why it shows judgment: - Not treating all data access the same (some access is read-only without consequences; some read requires approval before responding) - Recognizing that compliance changes the workflow, not just access - Decision: Different approval thresholds for different answer categories

Result: Built tier-specific workflows instead of assuming one-size-fits-all approval.

Question 4: Draft Scope for High-Risk Categories

Question: What should the AI definitely NOT try to answer, even with human review?

Why This Matters: - Tempting assumption: "AI can draft responses to anything if a human reviews them" - Reality: Some categories are so high-risk (or so outside AI's capability) that it shouldn't draft at all - Example: billing disputes (refund decisions), compliance questions (HIPAA-adjacent), customer escalations saying "I want to talk to a manager"

Changed thinking from: - "AI tries to draft everything; humans filter what's too risky" (reactive)

To: - "Define the high-risk categories upfront. AI learns not to touch those. For acceptable categories, AI drafts + human reviews. For high-risk, escalate immediately to a human who drafts."

Why it shows judgment: - Recognizing that even with perfect human review, some AI drafts are worse than useless (they burn time, not save it) - Refusing to overclaim AI's capability ("AI can draft anything with review" is false) - Decision: Explicit category-based guardrails, not just "review will catch everything"

Result: Avoided the confidence trap of "if humans review it, any AI output is OK." Built real constraints.

Question 5: Knowledge Base Health

Question: The knowledge base is supposed to be the source of truth. Is it? When was it last updated?

Why This Matters: - Scenario: Product changed a feature 6 months ago. Docs weren't updated. AI gets trained on old docs. AI tells customers outdated information, reps contradict it, customers get

confused. - The real question: "Can we trust the knowledge base enough to let AI cite it?"

Changed thinking from: - "We'll build the AI on our knowledge base" (trusting)

To: - "The knowledge base is 40% outdated. High-priority sections are fresh; low-priority sections haven't been touched in 8 months. We can't let the AI cite it broadly. Option A: Fix the KB first (3-week project). Option B: Scope the AI to only recent/popular articles and have it refuse others. Option C: Have human spot-check all AI citations."

Why it shows judgment: - Not assuming documentation is current (it almost never is) - Making a conscious tradeoff: delay launch to fix KB, or launch with narrow scope - Decision: Launch narrow scope + human spot-checks while KB cleanup happens in parallel

Result: Avoided confidently spreading stale information, and built in a mitigation (spot-checks) while the real fix (KB maintenance) happens separately.

What These 5 Questions Have in Common

1. **They don't accept assumptions at face value.** Each one challenges something the brief seemed to say was settled.
2. **They dig into the gap between theory and practice.** What does the process look like on paper vs. what do people actually do?
3. **They identify constraint-inducing facts.** Each answer becomes a hard constraint on what the AI can actually do.
4. **They show restraint.** Not "What should the AI do?" but "What should the AI NOT do? What can we not do safely?"

5. **They led to different recommendations.** The AI ended up narrower (Tier 1 troubleshooting only) and more careful (explicit guardrails) than the initial proposal.
-

Why These Matter for Interviews

If asked "What would you change about Meridian?", you can say:

"In discovery, I focused on the gap between documented process and what reps actually do. That's where I found the real constraints. For example, billing access turned out to require different approval workflows than troubleshooting, so we couldn't have a one-size-fits-all 'AI draft + human review' model. That changed how we scoped the feature."

This shows: - You ask hard questions during discovery (not just gathering requirements) - You reason about constraints (not just capabilities) - You distinguish different types of work (billing vs. troubleshooting aren't interchangeable) - You're willing to narrow scope when a clearer picture emerges

How to Defend This From Memory

If asked "Walk me through your discovery process":

1. Start with the original brief (narrow: "troubleshooting assistant")
2. Pick one constraint that surprised you (billing approval was the clearest one)
3. Explain the before/after (before: "AI can answer billing if a human reviews it" → after: "Billing requires Tier 2 approval gate, which is different from regular review")
4. End with the implication (scope changed, workflows changed)

You don't need to remember all 5 questions perfectly. You need to remember: - That you found hidden requirements through discovery -

That these requirements changed the design - One concrete example of a constraint you found

The Meta-Insight Worth Mentioning

"Discovery for AI-assisted work is harder than for typical software because you're not automating something that's already fully defined. You're augmenting human judgment. So you have to understand the actual judgment call—not just the happy path in the documentation—before you can design the right AI boundaries."

That's the actual lesson, and it's defensible from first principles, not memory.